

---

# CS6240

## Final Project

Purushottam Jha

Rushil Patel

Tushar Bhandari

Nishant Agarwal

---

# DATA

---

**Name:** Airline On-Time Performance and Causes of Flight Delays

**Size:** 4.34GB with 4336074861 records.

**Fields:** 55 Columns (such as FlightNo, Date, DepDelayMins, etc.)

**View of sample dataset:** (with few columns)

flightdate	flightnum	origin	origincityname	dest	destcityname	deptime	arrtime	arrdelayminutes
2007-10-02	415	IND	Indianapolis, IN	MDW	Chicago, IL	1102	1100	30.00
2007-10-02	917	IND	Indianapolis, IN	MDW	Chicago, IL	0728	0712	0.00
2007-10-02	1022	IND	Indianapolis, IN	MDW	Chicago, IL	1855	1842	0.00
2007-10-02	2081	IND	Indianapolis, IN	MDW	Chicago, IL	1530	1522	0.00
2007-10-02	1033	IND	Indianapolis, IN	PHX	Phoenix, AZ	1500	1544	0.00

Link to the dataset: <https://explore.data.gov/Transportation/Airline-On-Time-Performance-and-Causes-of-Flight-D/ar4r-an9z>

---

# TASKS

---

Task 1	Plain MapReduce	Finding Hubs and Spokes.
Task 2	HBase	a) On Time Arrival Performance of Airlines b) Average delay at each airport.
Task 3	Plain MapReduce	Calculate PageRank of each Airport.
Task 4	Hive and Pig Latin	Three legged round trip flights from Boston.

---

# TASK 1 : Hubs and Spokes

---

- Understanding HITS Algorithm
    - Iterative - 30 iterations [ 5- Small Machines : 1hr 8mins ].
    - Each Node [ Airport ] has HubValue and SpokeValue.
    - Normalization is done at each airport as follows
      - $$NewHubVal = HubVal / \sqrt{\sum (HubVals \text{ at Each Airport})^2}$$
  - 2 MR Jobs:
    - Convert file Data into Node Structure.
    - Perform HITS Algorithm iteratively.
-

# HITS in MapReduce

---

Job 2:

## **MAPPER:**

map(Key k, Value v) // Here Key: Object Value: Text

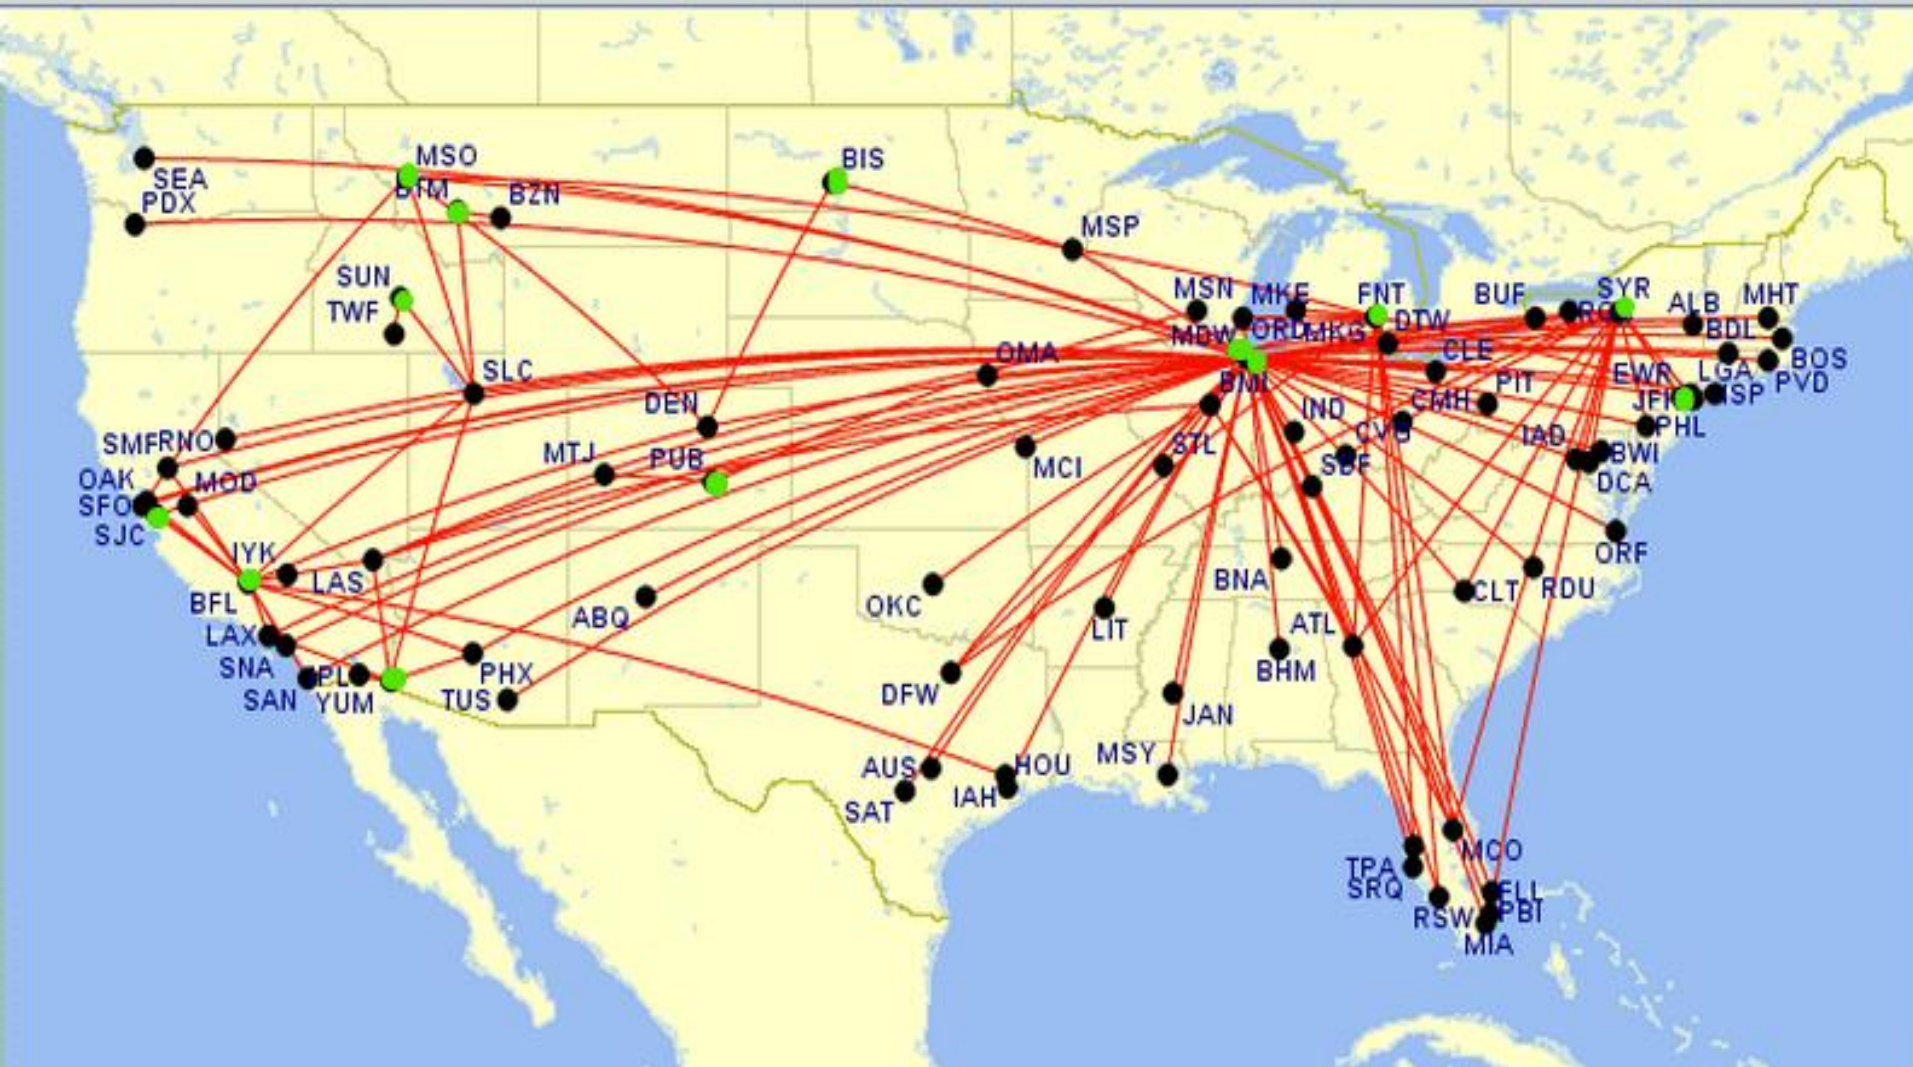
```
{
    Node n = parse(v)
    emit(n.nid, N )
    for(Node p in n.getInList()){
        n.setIsIn("YES")
        emit(p.nid,n)
    }
    for(Node p in n.getOutList()){
        n.setIsIn("NO")
        emit(p.nid,n)
    }
}
```

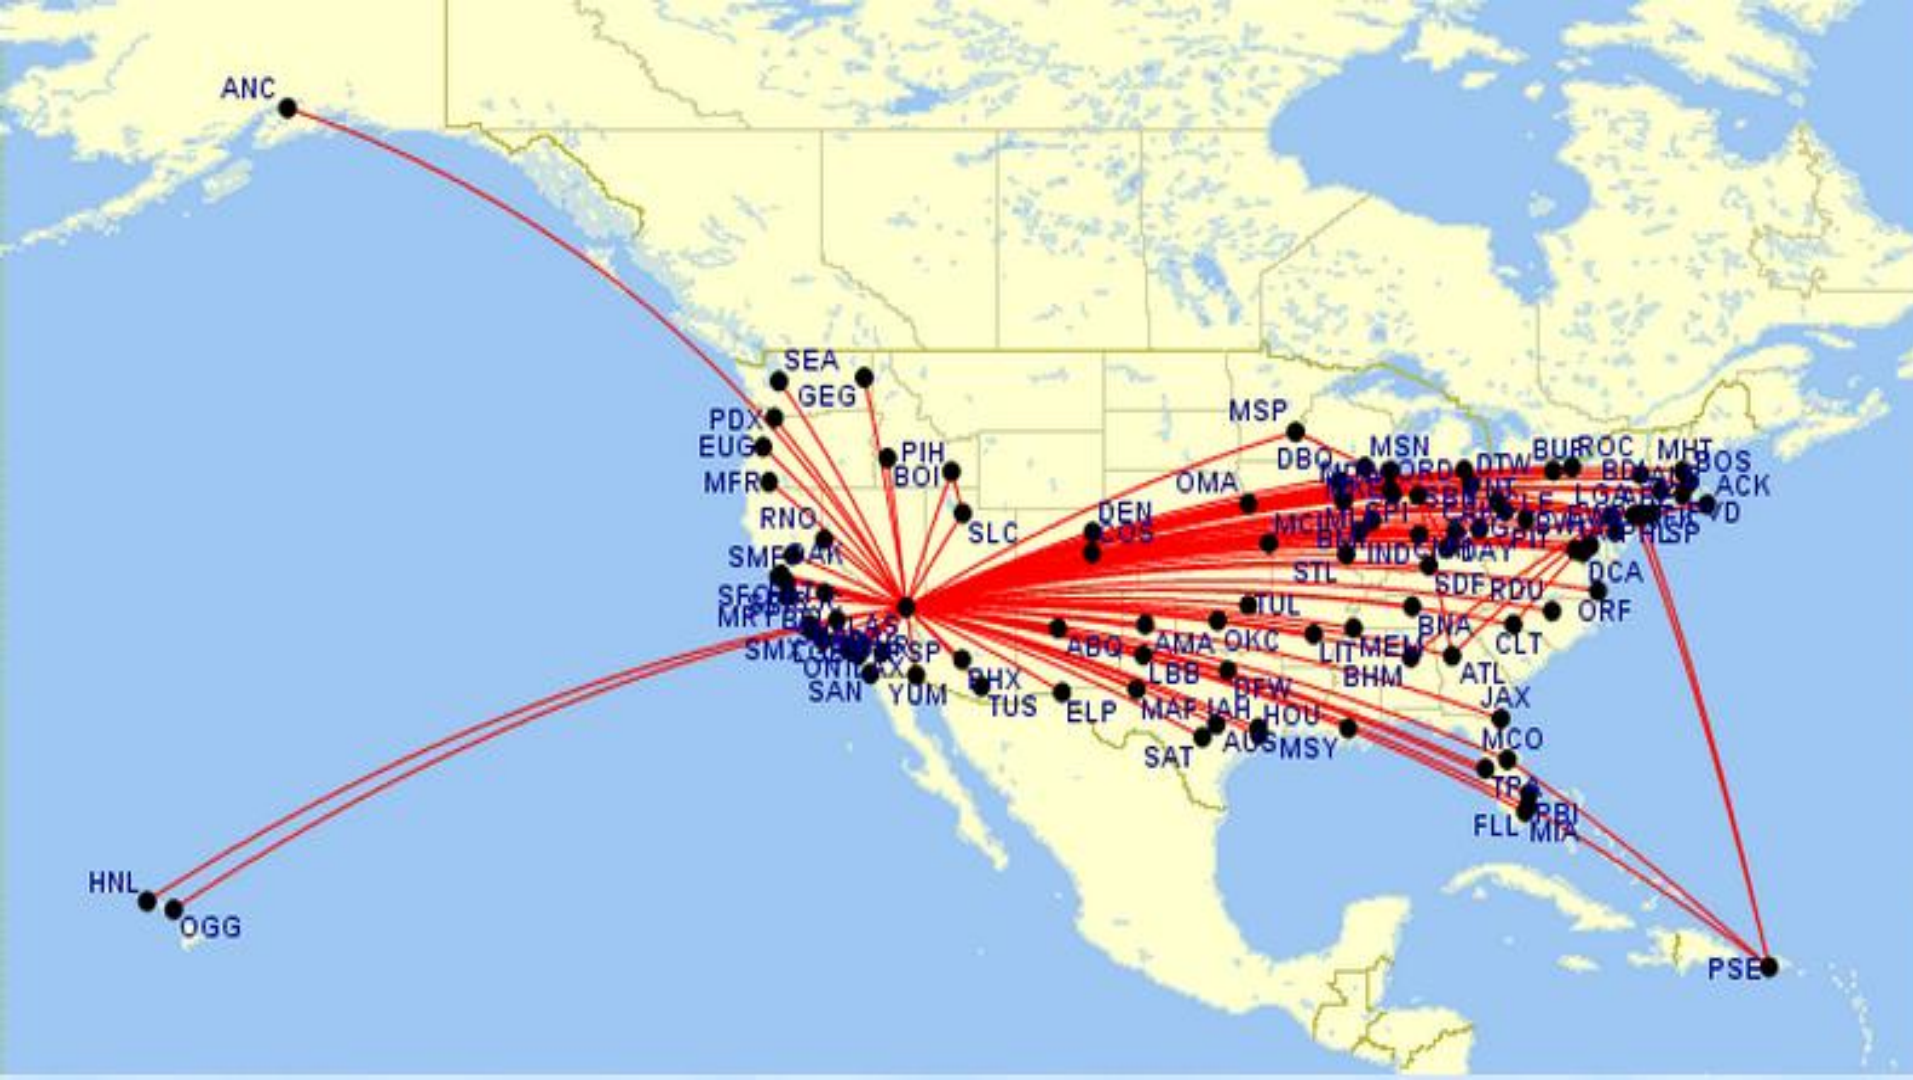
## **REDUCER**

```
setup()
{
    map = new Map(); // initialise hashmap
}
```

reduce(Key k, Values[v1,v2,v3,... ]) // Key : Text Values:[ Node n1. Node n2,.....]

```
{
    Node M = null
    HubVal = 0
    SpokeVal = 0
    for(Node n in values)
    {
        if(isNode(N))
            M=n
        else
            if(n.isInList())
                SpokeVal +=n.getHubVal()
            else
                HubVal += n.getSpokeVal()
    }
    map.put(n)
}
cleanup()
{
    for(val in map)
        HubNorm += Math.pow(val.Hubval,2)
        SpokeNorm +=Math.pow(val.SpokeVal,2)
    for(val in map)
        val.HubVal = val.HubVal/HubNorm
        val.SpokeVal = val.getSpokeVal/SpokeNorm
        emit(val,"")
}
```







# TASK 2 : HBase

---

## 1. On time arrival performance of airlines.

### MR Job 1:

```
// key = line offset , value = each record as String
map(key, value) {
    airlineId = value.getAirlineID();
    arrDelayMins = value.getArrDelayMinutes();
    emit(airlineId, arrDelayMins);
}

reduce(key, List[values]) {
    sum, total = 0,
    for delayValue in values:
        total ++;
        sum += delayValue;
    //compute average and emit in reducer
    averageDelayMins =(sum / total);
    // HBaseConnection is a utility class
    HBaseConnection.addRecord(tablename,
    averageDelayMins , "airlineID", "", key.toString());
}
```

### MR Job 2: (Map-Only)

```
// key is averageDelay
// value is airlineID
map( key, value) {

    //convert average delay
    // from Bytes to DoubleWritable
    averageDelay = row.get()

    // get airline ID from result
    airlineID = value.getValue(CF, ATTR)

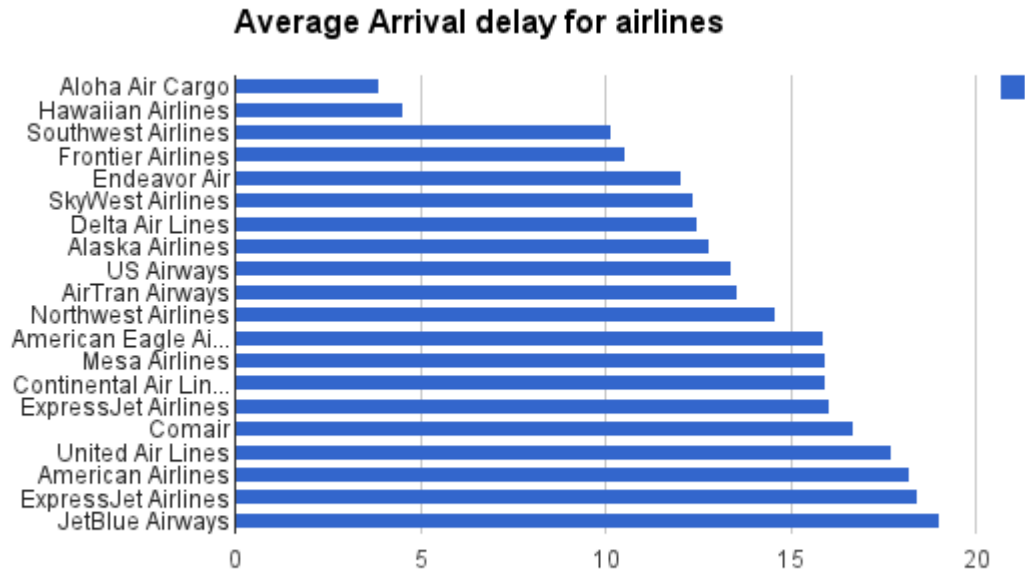
    emit(averageArrivalDelay, airlineID);
}
```



# HBase

---

## 1. On-time arrival performance of airlines.



# HBase

---

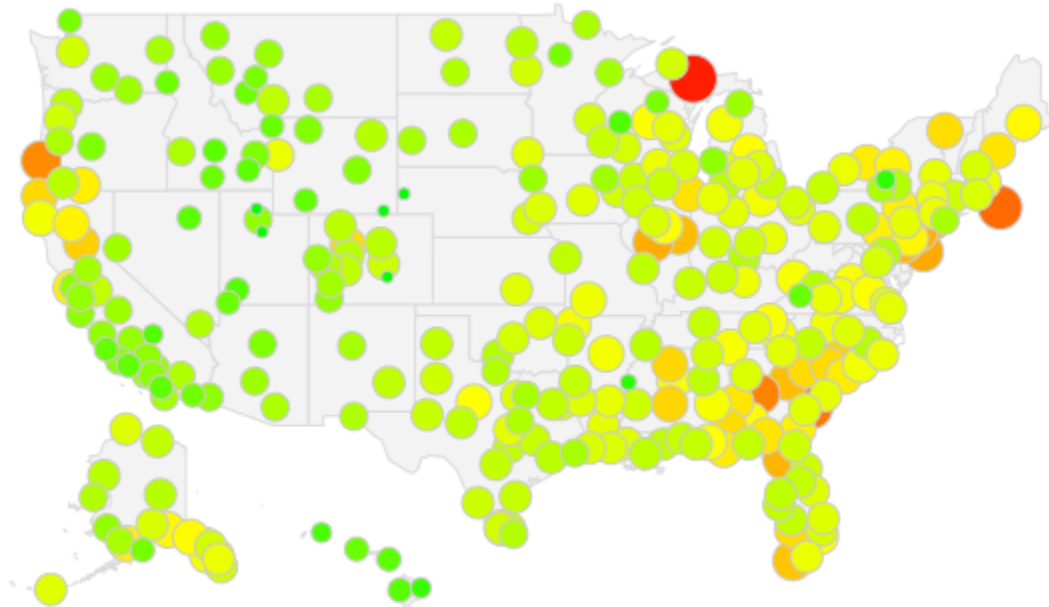
## Performance numbers on EMR

Program	No. of Worker Machines	Machine Type	Time Taken
HBase1 (Task 1)	2	large	8 minutes
HBase 1 (Task 1)	2	large	9 minutes
HBase2 (Task 2)	5	large	4 minutes
HBase2 (Task 2)	5	large	4 minutes

# HBase

---

## 2. Average delay at each airport



# Task 3: Pagerank

---

- Multiple flights having same source and destination does affect the pagerank value
  - However, the data transfer can be minimized by using in-mapper aggregation and/or combiner
  - This way if there are 10 flights from BOS -> JFK, instead of emitting (JFK, [BOS,BOS,BOS,BOS.....]),  
emit(JFK,[(BOS,10),(MIA,3),(SEA,5)])
-

# Pagerank discrepancy

---

CLE and LAX switch their position in the order

	Without count	With count
<b>BOS</b>	0.03509803921568627	0.03509803921568627
<b>CLE</b>	0.05751257812154845	0.06333333333333332
<b>ORD</b>	0.03509803921568627	0.03509803921568627
<b>LAX</b>	0.06777777777777777	0.03333333333333333

# Aggregation Vs No-Aggregation

---

It took double the resources to achieve the same runtime

	Aggregation	No-Aggregation
# of machines	10	20
type	small	small
map output records	163749	26790152
time taken (secs)	295	297

# Same data, more workers

---

More workers can sometimes decrease the performance

<b>Pagerank</b>	<b>10 machines</b>	<b>20 machines</b>
<b>Iteration #1</b>	140	170
<b>Iteration #2</b>	143	174
<b>Iteration #3</b>	141	172
<b>Iteration #4</b>	142	171



# TASK 4 : Hive and PigLatin

---

Shortest three-legged round trip journey.

Hive Code:

```
INSERT OVERWRITE DIRECTORY 's3n://finalproj-puru/hiveoutput'
Select a.origin, a.dest, b.dest, c.dest, a.flightdate, a.DepTime, a.ArrTime, b.DepTime,
b.ArrTime, c.DepTime, c.ArrTime, (a.actualelapsedtime + b.actualelapsedtime + c.
actualelapsedtime + (b.DepTime-a.ArrTime) + (c.DepTime-b.ArrTime)) AS TotalTime from
flight a JOIN flight b on (a.dest = b.origin) JOIN flight c on (b.dest = c.origin) where
a.flightdate = '2008-01-01' AND b.flightdate = '2008-01-01' AND c.flightdate = '2008-01-
01' AND a.origin = 'BOS' AND c.dest = 'BOS'
AND a.DepTime < a.ArrTime
AND b.DepTime < b.ArrTime
AND c.DepTime < c.ArrTime
AND b.DepTime > a.ArrTime
AND c.DepTime > b.ArrTime AND b.DepTime - a.ArrTime > 100 AND c.DepTime - b.ArrTime >
100
AND a.cancelled != 1 AND b.cancelled != 1 AND c.cancelled != 1
Order by TotalTime LIMIT 20;
```

# Hive and PigLatin

---

## PigLatin Code:

```
Flights1_Data = FILTER Flights1_Data by (orig1 == 'BOS') AND (flightDate1 == '2008-01-01') AND (cancelled1 != 1);
Flights2_Data = FILTER Flights2_Data by (flightDate2 == '2008-01-01') AND (cancelled2 != 1);
Flights3_Data = FILTER Flights3_Data by (dest3 == 'BOS') AND (flightDate3 == '2008-01-01') AND (cancelled3 != 1);

f1f2 = JOIN Flights1_Data BY (dest1), Flights2_Data BY (orig2);
f1f2 = FILTER f1f2 BY depTime2 > arrTime1;

f1f2f3 = JOIN f1f2 BY (dest2), Flights3_Data BY (orig3);
f1f2f3 = FILTER f1f2f3 BY depTime3 > arrTime2;

f1f2f3 = FILTER f1f2f3 BY ((depTime2-arrTime1) > 100) AND ((depTime3-arrTime2) > 100);

final = FOREACH f1f2f3 GENERATE orig1,dest1,dest2,dest3,flightDate1,depTime1,arrTime1,depTime2,arrTime2,depTime3,
arrTime3, (actualElapsedTime1 + actualElapsedTime2 + actualElapsedTime3 + (depTime2 - arrTime1) + (depTime3 -
arrTime2)) AS totalTripTime;

final = ORDER final BY totalTripTime;
final = limit final 20;
STORE final INTO '$OUTPUT';
```

# Hive and PigLatin

---

Origin	Dest 1	Dest 2	End	Flight Date	F1: DepTime	F1: ArrTime	F2: DepTime	F2: ArrTime	F3: DepTime	F3: ArrTime	TotalTripTime
BOS	DCA	LGA	BOS	1/1/2008	847	1028	1131	1223	1328	1435	428
BOS	DCA	LGA	BOS	1/1/2008	847	1028	1131	1223	1339	1442	435
BOS	DCA	LGA	BOS	1/1/2008	1602	1735	1852	2004	2107	2200	438
BOS	IAD	LGA	BOS	1/1/2008	1021	1211	1320	1420	1528	1625	444
BOS	DCA	LGA	BOS	1/1/2008	847	1028	1131	1223	1354	1452	445
BOS	JFK	PHL	BOS	1/1/2008	1434	1542	1645	1832	1938	2041	447
BOS	LGA	DCA	BOS	1/1/2008	1128	1229	1352	1502	1619	1738	450
BOS	IAD	LGA	BOS	1/1/2008	1021	1211	1320	1420	1530	1632	451
BOS	IAD	LGA	BOS	1/1/2008	1021	1211	1320	1420	1556	1647	466

# Hive and PigLatin

---

## Performance Numbers on EMR

Program	No. of Machines	Machine Type	Time Taken
HIVE	11	small	32 Mins
HIVE	20	small	15 Mins
PigLatin	11	small	14 Mins 49 Secs
PigLatin	20	small	13 Mins 23 Secs

---

**THANK YOU**

---